# Parth Gawande

San Jose, CA | parthgawande2000@gmail.com | 585-625-9131 | linkedin.com/in/parthgawande | github.com/parthgawande

## SUMMARY

Data Engineer with 3+ years of experience building batch and real-time pipelines, data models, and analytics infrastructure for Operations, Finance, and leadership teams. Works hands-on across Palantir Foundry, AWS, Python, SQL, and Spark, from raw ingestion through data quality checks to Power BI and Grafana reporting. Owns pipelines end-to-end, from design through production support.

## SKILLS

**Languages:** Python (Pandas, NumPy, PySpark), SQL (CTEs, Window Functions, Query Optimization), Bash
**Data Engineering:** ETL/ELT, Apache Spark, Apache Airflow, dbt, Kafka, Batch & Real-Time Processing, Data Pipeline Orchestration
**Cloud & Warehousing:** AWS (S3, Lambda, RDS, EC2, Glue), Snowflake, Databricks, Palantir Foundry & AIP
**Data Modeling & Quality:** Dimensional Modeling (Star Schema), Schema Design, Data Validation, Great Expectations, SLA Monitoring
**GenAI & LLMs:** AWS Bedrock, LangGraph, RAG, FAISS, Vector Embeddings, Guardrails, Palantir AIP, LLM Evaluation (DeepEval)
**DevOps & BI:** Git, CI/CD, Docker, Agile, Power BI, Tableau, Grafana, SAP

## EXPERIENCE

**Data Engineer** **March 2026 - October 2026**
*Pacific Gas & Electric* *San Jose, CA*

- Saved $3M+ and removed manual data prep for 8 teams by building a Palantir Foundry pipeline that processes 50M+ inspection records a year in SQL and Python, holding 99%+ consistency across compliance and field metrics.
- Accelerated Asset Strategy planning by 3 weeks and cut manual scope definition by 80% for 40+ inspectors by automating inspection cycle design with Foundry data models feeding live Power BI dashboards.
- Freed 14 hours/week of manual mapping for 40+ field inspectors by building an automated assignment system in Palantir Foundry that combines workload, location, and inspection data with Python routing logic.
- Built LLM-powered Palantir AIP workflows in Foundry that let inspection and compliance teams query pipeline data in plain language, reducing ad hoc report requests to the data team.
- Cut reporting turnaround by 75% across 5 verticals (SIC, SIQ, Training, Finance, Asset Strategy) by mapping manual steps with Lean principles, modeling each team's data in Foundry, and launching self-service dashboards.

**Data Engineer** **September 2025 - March 2026**
*Saayam For All* *San Jose, CA*

- Reduced matching latency by 25% and doubled daily match capacity across 200K+ volunteer profiles and 40K+ opportunities by rewriting the matching pipeline's SQL with window functions and tuning the matching algorithm.
- Improved matching accuracy by 30% by tracing a 45% volunteer drop-off in 3 programs to incomplete opportunity descriptions through cohort analysis, then sharing the fixes through Tableau dashboards.

**Data Engineer** **January 2021 - December 2022**
*StandardWings Technologies Pvt. Ltd.* *Nashik, India*

- Led a zero-downtime migration from SAP and legacy systems to AWS, building an S3, Lambda, and RDS layer for GPS and telemetry from 120+ trucks that reached a 99.2% record match in post-migration validation.
- Lowered missed collections by 18% and improved route efficiency by 15% over 6 months vs. the manual-scheduling baseline by building real-time GPS and sensor processing in Python on AWS that drove automated dispatch.
- Saved $400K a year in repair costs and cut incident discovery time by 72% vs. the manual-inspection baseline by deploying Python anomaly detection models on AWS sensor data across a 120+ vehicle fleet.
- Standardized metrics across 4 teams and removed 20-minute routing bottlenecks by defining 21 KPI specifications with stakeholders and instrumenting the AWS data layer behind self-service Grafana dashboards.

## PROJECTS

**BankGuard: Guardrailed GenAI Compliance Assistant**

- Blocked 98% of unsafe and off-policy prompts across 55 test cases by building a multi-agent banking assistant on AWS Bedrock (Claude) with Guardrails, Titan embeddings, FAISS RAG, PII masking, and CloudTrail audit logs.

**MemoryOS: Self-Correcting Multi-Agent Assistant with Long-Term Memory**

- Achieved 0.85 answer relevancy and 0.97 faithfulness on 33 DeepEval test cases by building a LangGraph planner-executor-critic agent with GPT-4o, mem0 persistent memory, and web search, code, and calendar tools.

## EDUCATION

**MS in Information Technology and Analytics** **January 2023 - May 2025**
*Rochester Institute of Technology | GPA: 3.9Rochester, New York*
**B.Tech in Computer Science and Engineering** **July 2018 - July 2022**
*MIT World Peace University | GPA: 3.8Pune, India*